

Meets Specifications

Dear Udacian, the project is very well implemented and meets the specifications! Congratulations on successfully completing the project. As a next step, please go through the resource: [human-like robot hand trained to manipulate physical objects with unprecedented dexterity](#).

All the best for the next project! 🙌

Training Code

| | |
|---|--|
| ✓ | The repository includes functional, well-documented, and organized code for training the agent. |
| | The repository contains jupyter notebook, code files, readme, and project report. The code is functional. Great job solving the environment using Deep Deterministic Policy Gradients. |
| ✓ | The code is written in PyTorch and Python 3. |
| | <div>The code is written in pytorch and python3.<ul style="list-style-type: none">A good read: PyTorch vs TensorFlow—spotting the difference</div> |
| ✓ | The submission includes the saved model weights of the successful agent. |
| | Thanks for including the saved model weights of the successful agent. |

README

| | |
|---|---|
| ✓ | The GitHub submission includes a <code>README.md</code> file in the root of the repository. |
| ✓ | The README describes the the project environment details (i.e., the state and action spaces, and when the environment is considered solved). |
| | Awesome work providing the project environment details in the README. State space, action space, reward function and when the environment is considered solved is specified very informatively. |
| ✓ | The README has instructions for installing dependencies or downloading needed files. |
| | Proper instructions have been specified in the README to download the necessary files. |
| ✓ | The README describes how to run the code in the repository, to train the agent. For additional resources on creating READMEs or using Markdown, see here and here . |

Report

| | |
|---|---|
| ✓ | The submission includes a file in the root of the GitHub repository (one of <code>Report.md</code> , <code>Report.ipynb</code> , or <code>Report.pdf</code>) that provides a description of the implementation. |
| | Report has been included in the root of the github repository. |
| ✓ | The report clearly describes the learning algorithm, along with the chosen hyperparameters. It also describes the model architectures for any neural networks. |
| | <div>The report is rather informative providing an insight on every aspect of the project which includes implementation, model architectures, hyperparameters, rewards, future works.<ul style="list-style-type: none">Good choice to implement the DDPG algorithm.It is found to work very well with continuous action space.Good Implementation of the Actor and Critic networks.Good decision to use replay buffer to store and recall experience tuples.Good job using the target networks for Actor and Critic, as suggested in the original paper.Good choice to use tau to update the target network.</div> |
| ✓ | <div>A plot of rewards per episode is included to illustrate that either:<ul style="list-style-type: none"><i>[version 1]</i> the agent receives an average reward (over 100 episodes) of at least +30, or<i>[version 2]</i> the agent is able to receive an average reward (over 100 episodes, and over all 20 agents) of at least +30.The submission reports the number of episodes needed to solve the environment.</div> |
| | <div><h3>Awesome</h3><ul style="list-style-type: none">The agent seems to perform very well!The agent is able to achieve an average score of 30+ over last 100 episodes in just 174 episodes!The submission discusses the rewards plot obtained clearly.</div> |

| | |
|---|--|
| ✓ | The submission has concrete future ideas for improving the agent's performance. |
| | <div>Great job providing the ideas to experiment more in future with the project!</div> <div>You should try implementing Prioritized Experience Replay also. It helps to improve the performance and significantly reduces the training time. A fast implementation of Prioritized Experience Replay is possible using a special data structure called Sum Tree. I found a good implementation here.</div> <div>Also, I request you to check the following posts to get familiar with more reinforcement learning algorithms.<ul style="list-style-type: none">Asynchronous Actor-Critic Agents (A3C)Trust Region Policy Optimization (TRPO) and Proximal Policy Optimization (PPO)</div> <div>As pointed in the report, you should definitely try solving the environment using Proximal Policy Optimization algorithm. For reference, you may check this github repository.</div> |